

# CNN

February 2, 2018

## 1 Convolutional Neural Network

### 1.1 Import Dependencies

```
In [1]: %matplotlib inline
        from imp import reload

        import itertools
        import numpy as np
        import utils; reload(utils)

        from utils import *
        from __future__ import print_function
        from sklearn.metrics import confusion_matrix, classification_report, f1_score
        import matplotlib.pyplot as plt

In [2]: from keras.preprocessing import sequence
        from keras.models import Sequential
        from keras.layers import Dense, Dropout, Activation
        from keras.layers import Embedding, SpatialDropout1D
        from keras.layers import LSTM
        from keras.layers import Conv1D, GlobalMaxPooling1D
        from keras.layers import Flatten
        from keras.datasets import imdb
        from keras.utils import plot_model
        from keras.utils.vis_utils import model_to_dot

        from IPython.display import SVG
        from IPython.display import Image
```

```
/home/konstantin/.local/lib/python3.5/site-packages/h5py/__init__.py:36: FutureWarning: Conversion
  from the numpy float64 dtype to the numpy float32 dtype is deprecated
from ..conv import register_converters as _register_converters
Using TensorFlow backend.
```

### 1.2 Configure Parameters

```
In [3]: # Embedding
        embedding_size = 50
```

```

max_features = 5000
maxlen = 400

# Convolution
kernel_size = 3
pool_size = 4
filters = 250

# Dense
hidden_dims = 250

# Training
batch_size = 64
epochs = 4

```

### 1.3 Data Preparation

```
In [4]: (x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)
```

```
In [5]: # Pad sequences
x_train = sequence.pad_sequences(x_train, maxlen=maxlen)
x_test = sequence.pad_sequences(x_test, maxlen=maxlen)

print('Train data size:', x_train.shape)
print('Test data size:', x_test.shape)
```

```

Train data size: (25000, 400)
Test data size: (25000, 400)

```

### 1.4 Modelling

```
In [6]: model = Sequential()

# we start off with an efficient embedding layer which maps
# our vocab indices into embedding_dims dimensions
model.add(Embedding(max_features,
                    embedding_size,
                    input_length=maxlen))
model.add(Dropout(0.2))

model.add(Conv1D(filters,
                kernel_size,
                padding='valid',
                activation='relu',
                strides=1))
model.add(GlobalMaxPooling1D())
```

```

# We add a vanilla hidden layer:
model.add(Dense(hidden_dims))
model.add(Dropout(0.2))
model.add(Activation('relu'))

# We project onto a single unit output layer, and squash it with a sigmoid:
model.add(Dense(1))
model.add(Activation('sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

model.summary()

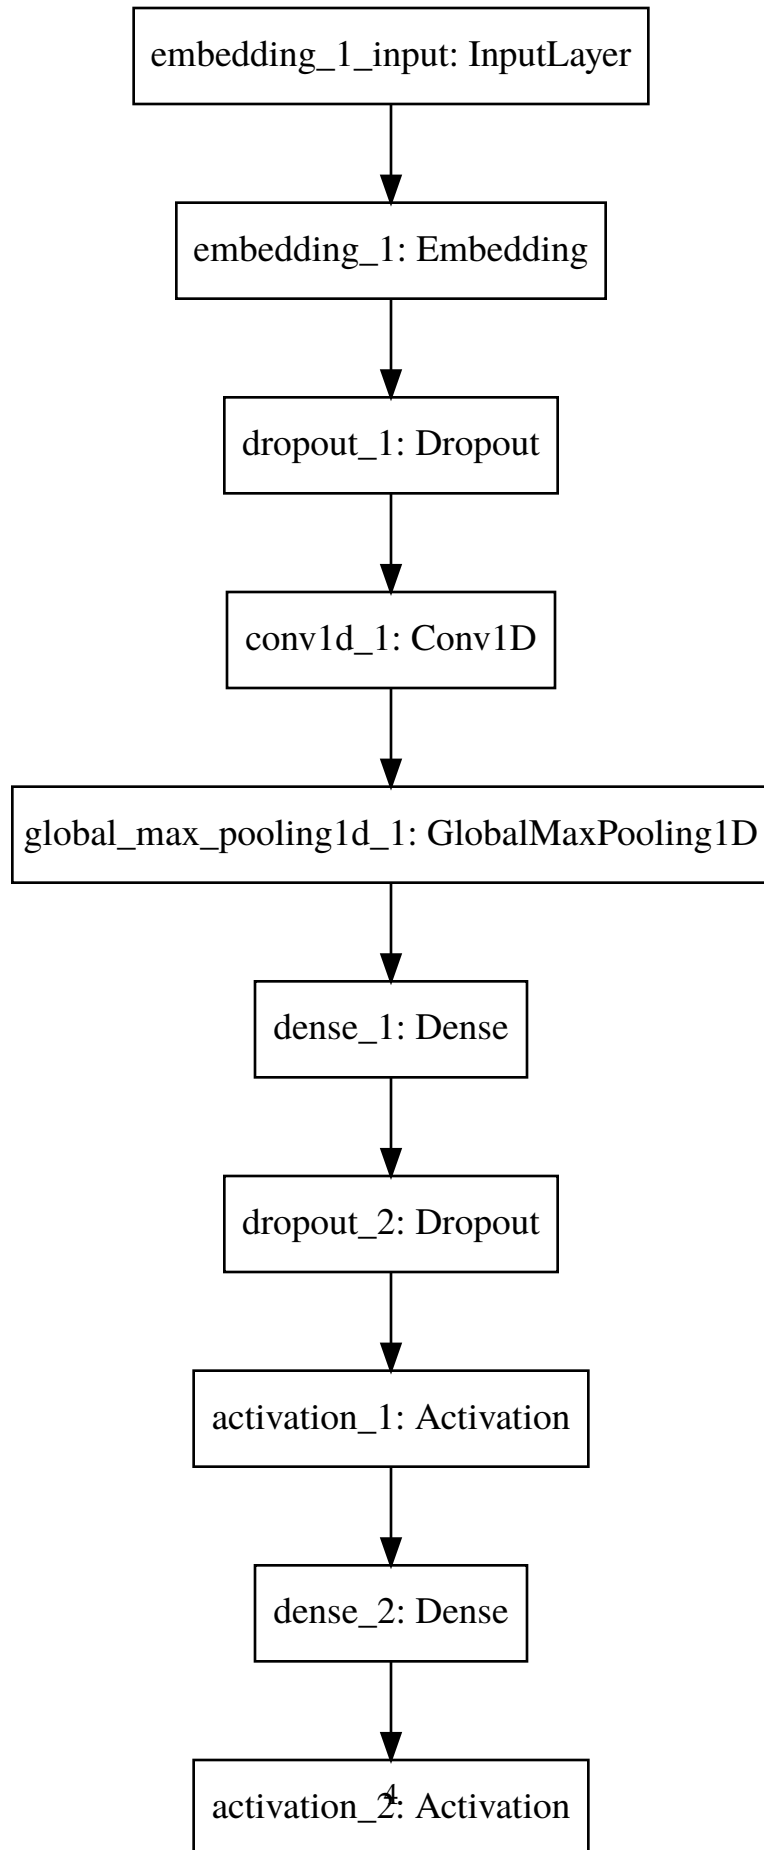
plot_model(model, to_file='model_cnn.png', show_shapes=True)
Image(filename = 'model_cnn.png')
SVG(model_to_dot(model).create(prog='dot', format='svg'))

```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 400, 50)	250000
dropout_1 (Dropout)	(None, 400, 50)	0
conv1d_1 (Conv1D)	(None, 398, 250)	37750
global_max_pooling1d_1 (Glob	(None, 250)	0
dense_1 (Dense)	(None, 250)	62750
dropout_2 (Dropout)	(None, 250)	0
activation_1 (Activation)	(None, 250)	0
dense_2 (Dense)	(None, 1)	251
activation_2 (Activation)	(None, 1)	0

Total params: 350,751  
 Trainable params: 350,751  
 Non-trainable params: 0

Out[6]:



## 1.5 Evaluation

In [7]: *# Train the model*

```
model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          validation_data=(x_test, y_test),
          verbose=1)
```

Train on 25000 samples, validate on 25000 samples

Epoch 1/4

25000/25000 [=====] - 6s 247us/step - loss: 0.4395 - acc: 0.7752 - val\_

Epoch 2/4

25000/25000 [=====] - 5s 209us/step - loss: 0.2401 - acc: 0.9019 - val\_

Epoch 3/4

25000/25000 [=====] - 5s 205us/step - loss: 0.1675 - acc: 0.9370 - val\_

Epoch 4/4

25000/25000 [=====] - 5s 206us/step - loss: 0.1167 - acc: 0.9590 - val\_

Out[7]: <keras.callbacks.History at 0x7fcd5141588>

In [8]: *# Evaluate model*

```
score, acc = model.evaluate(x_test, y_test, batch_size=batch_size)
preds = model.predict_classes(x_test, batch_size=batch_size)
```

25000/25000 [=====] - 1s 41us/step

In [9]: *# Save the model weights*

```
model_path = 'models/'
model.save(model_path + 'cnn_model.h5')
model.save_weights(model_path + 'cnn_weights.h5')
```

In [10]: *def plot\_confusion\_matrix*(cm, classes,

```
    normalize=False,
    title='Confusion matrix',
    cmap=plt.cm.Blues):
```

```
    """
```

```
    This function prints and plots the confusion matrix.
```

```
    Normalization can be applied by setting `normalize=True`.
```

```
    """
```

```
    if normalize:
```

```
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
```

```
        print("Normalized confusion matrix")
```

```
    else:
```

```

        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)


    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

In [11]: # Confusion Matrix
         cm = confusion_matrix(y_test, preds)
         plot_confusion_matrix(cm, {'negative': 0, 'positive': 1})

Confusion matrix, without normalization
[[10707  1793]
 [ 1020 11480]]

```



```
In [12]: # F1 score
f1_macro = f1_score(y_test, preds, average='macro')
f1_micro = f1_score(y_test, preds, average='micro')

print('Test accuracy:', acc)
print('Test score (loss):', score)
print('')
print('F1 Score (Macro):', f1_macro)
print('F1 Score (Micro):', f1_micro)
```

Test accuracy: 0.887479999961853

Test score (loss): 0.30362622849464416

F1 Score (Macro): 0.8873723227145909

F1 Score (Micro): 0.88748